

# CSC490 Assignment 5 - Deleting Prod

Team:

Son Nguyen (ID 1009656560), Kyle (ID 1007785229),

Daniel (ID 1008035378), Jinbo (ID 1004821419)

Project: ArXplorer - Semantic Search for arXiv

## Part One

Our goals were:

- Measure where our system actually spends time.
- Identify the main bottlenecks in the embedding and search pipeline.
- Propose concrete optimizations we can implement in later parts or future work.

We focused on the embedding, indexing, and search components, since these are the core of our LLM-based semantic search system.

Profiling was done with Python 3.10 on Windows using the script:

- `a5/profile_a5_part1.py`

Each profiling run produced a “.prof” file, which we saved under:

- `a5/profiling_results/`
- 

### 2. Functions Profiled

---

We profiled the following five functions:

1. `_encode_text()`
  - Module: `src.core.pipeline.EmbeddingGenerator`
  - Reason: Lowest-level “text → embedding” operation, used throughout the system.
2. `generate_embeddings()`
  - Module: `src.core.pipeline.EmbeddingGenerator`
  - Reason: Generates full embeddings for a `ProcessedPaper` (title + abstract, etc.). Called once per paper.
3. `build_index()`
  - Module: `src.core.pipeline.VectorIndexer`
  - Reason: Builds the FAISS vector index that powers semantic search.
4. `search_papers_by_text()` (mocked)
  - For Part 1, implemented via a `FakeRepo` inside `profile_a5_part1.py`.
  - Reason: Represents the search-by-text API path without requiring a real MongoDB setup.
5. `get_papers_by_categories()` (mocked)
  - Also implemented via `FakeRepo`.
  - Reason: Represents category-based search latency and async overhead.

For the last two functions we intentionally used a fake repository instead of a real MongoDB client. The goal in Part 1 is to profile code paths and runtime behavior, not to debug database configuration.

---

### 3. Profiling Setup

---

We used Python's cProfile module. The core helper in profile\_a5\_part1.py looks like this:

Function run\_profile(label, func\_str):

- Prints a header for the profiling section.
- Builds a filename such as "profile\_encode\_text\_YYYYMMDD\_HHMMSS.prof".
- Calls cProfile.run(func\_str, filename).
- Loads the profile file into pstats.Stats and prints the top 20 entries, sorted by total time (totime).

Example usage:

- run\_profile("encode\_text", "profile\_encode\_text()")
- run\_profile("generate\_embeddings", "profile\_generate\_embeddings()")
- run\_profile("build\_index", "profile\_build\_index()")
- run\_profile("search\_text", "asyncio.run(profile\_search\_by\_text())")
- run\_profile("search\_category", "asyncio.run(profile\_search\_by\_category())")

We kept only the latest profile file per function (to avoid clutter) and used the console output (top 20 lines sorted by totime) for our analysis.

---

#### 4. Profiling Results (Per Function)

---

##### 4.1 \_encode\_text()

Profile file: profile\_encode\_text\_... .prof

Total runtime (latest run): approximately 0.73 seconds

Role:

- Takes a single string (e.g., title or abstract).
- Tokenizes it with a HuggingFace tokenizer.
- Runs a transformer model (sentence-transformers/all-MiniLM-L6-v2).
- Returns one embedding vector.

Top bottlenecks (by totime):

- SSL socket reads (method 'read' of '\_ssl.\_SSLSocket'): about 0.17 seconds on the measured run.
- File I/O for tokenizer and model files (nt.stat, io.open, io.open\_code).
- torch.\_C.\_nn.linear (transformer dense layers).
- JSON parsing for model configuration.

Interpretation:

- The first run is dominated by model and tokenizer loading (and sometimes remote cache access via SSL).

- After the model is fully loaded and cached, subsequent calls are much faster and dominated by transformer forward pass.
  - This function is fundamental; any optimization here improves all higher-level embedding functions.
- 

#### 4.2 generate\_embeddings()

Profile file: profile\_generate\_embeddings\_... .prof  
Total runtime (latest run): approximately 0.35 seconds

Role:

- Takes a ProcessedPaper object containing:
  - cleaned\_title
  - cleaned\_abstract
  - extracted\_keywords, etc.
- Calls `_encode_text()` for the relevant fields (usually title + abstract).
- Produces a joint embedding representation for the paper.

Top bottlenecks:

- SSL reads (on some runs) and transformer model state loading (especially for the first call).
- `torch.C.nn.linear` and `torch.layer_norm`.
- A few JSON decode calls and model loading functions.

Interpretation:

- `generate_embeddings()` is essentially a small wrapper that calls `_encode_text()` multiple times.
- After warm-up, most cost comes from the transformer forward passes.
- Around 0.35 seconds per paper on CPU is acceptable for small datasets, but becomes significant for large-scale indexing.

Importance:

- This function is called for every paper we want to index.
  - It directly impacts offline preprocessing time and the index-building pipeline.
- 

#### 4.3 build\_index()

Profile file: profile\_build\_index\_... .prof  
Total runtime (latest run): approximately 2.09 seconds for a test batch of 50 papers.

Role:

- Generates dummy ProcessedPaper objects.
- Uses EmbeddingGenerator to create embeddings.
- Builds a FAISS index (e.g., FLAT index) over these embeddings.

Top bottlenecks:

- `torch._C._nn.linear`: about 0.81 seconds.
- `torch.layer_norm`: about 0.15 seconds.
- `torch._C._nn.gelu`: about 0.10 seconds.
- `torch._C._nn.scaled_dot_product_attention`.
- BERT model forward passes in `transformers (modeling_bert.py)`.

Interpretation:

- Most time is spent generating embeddings, not inside FAISS itself.
- FAISS index creation is relatively cheap for 50 vectors, but embedding generation dominates.
- As the number of papers grows, `build_index()` becomes a combination of “embedding many papers” plus “FAISS indexing”.

Importance:

- Determines how fast we can build or rebuild our semantic index.
- For production, this would be part of an offline pipeline that we want to make efficient.

---

#### 4.4 `search_papers_by_text()` (Mocked)

Profile file: `profile_search_text_... .prof`

Total runtime (latest run): approximately 0.002 seconds

Role in Part 1:

- In `profile_a5_part1.py`, we define a `FakeRepo` with an `async` method `search_papers_by_text`.
- The method simply returns a list of fake documents with a given limit (e.g., 10 items).
- We then profile `asyncio.run(profile_search_by_text())`.

Key observations from the profile:

- Almost all time is spent in `asyncio` event loop functions:
  - `base_events.run_until_complete`
  - `windows_events._poll`
  - `socketpair` and overlapped I/O operations.
- Our own application logic (creating a list of dictionaries) is negligible.

Interpretation:

- This setup measures the overhead of the `async` infrastructure rather than real DB latency.
- In a real deployment with MongoDB, the major cost would come from:
  - Network I/O to MongoDB.
  - Query execution and index usage.

Reason for mocking in Part 1:

- Avoids failure due to missing MongoDB collections or text indexes.
  - Keeps the focus of Part 1 on profiling Python code, not database configuration.
- 

#### 4.5 get\_papers\_by\_categories() (Mocked)

Profile file: profile\_search\_category\_... .prof

Total runtime (latest run): approximately 0.003 seconds

Role in Part 1:

- Similar to search\_papers\_by\_text(), we use a FakeRepo with get\_papers\_by\_categories.
- The mock method returns a list of fake documents where “category” is set to the first requested category.

Key observations:

- Again, most of the time shows up inside asyncio event loop internals.
- The user-level code is almost free compared to the async overhead.

Interpretation:

- As with the text search case, the real bottleneck in a production environment would be MongoDB queries and indexes, not the wrapper function.
- 

### 5. Cross-cutting Observations

---

#### 1. Model loading vs. steady-state performance

- The first call to embedding functions (\_encode\_text and generate\_embeddings) is relatively slow due to:
  - Download and caching of the HuggingFace model.
  - Tokenizer and configuration file loading.

- After warm-up, runtime is dominated by the transformer forward pass in PyTorch.

#### 2. Embedding computation dominates index building

- In our build\_index test, the bulk of the time is spent calling generate\_embeddings for 50 papers.
- FAISS index construction itself is relatively cheap at this scale.

#### 3. Async search overhead is small

- With the FakeRepo, search\_papers\_by\_text and get\_papers\_by\_categories complete in a few milliseconds.
  - The overhead is almost entirely the asyncio event loop.
  - Once we incorporate MongoDB, the main additional cost will be round trips and query execution.
- 

### 6. Optimization Ideas

---

Based on the profiling results, we identified several optimizations:

## 6.1 Cache tokenizer and model

Problem:

- If EmbeddingGenerator re-initializes models or tokenizers frequently, we pay repeated initialization costs.

Idea:

- Ensure the tokenizer and model are loaded once in EmbeddingGenerator.**init** and reused across all calls.
- Optionally, treat EmbeddingGenerator as a singleton in the application layer so the model is only created once per process.

Expected benefit:

- Removes expensive re-initialization on repeated requests.
  - Reduces cold-start overhead in production.
- 

## 6.2 Batch embedding for title and abstract

Problem:

- generate\_embeddings calls \_encode\_text multiple times for different fields of the paper (title, abstract, etc.).

Idea:

- Use batch encoding with the underlying model (e.g., encode [title, abstract] in one call).
- For bulk operations like indexing, also process multiple papers per batch.

Expected benefit:

- Better CPU/GPU utilization and less Python overhead.
  - Potentially 1.5x or better speedup for per-paper embedding generation.
- 

## 6.3 Avoid rebuilding FAISS index from scratch

Problem:

- build\_index currently assumes we generate embeddings and build the FAISS index every time we call it.

Idea:

- Move embedding generation into a separate offline preprocessing step.
- Save the FAISS index once and reload it on startup using:
  - faiss.write\_index
  - faiss.read\_index

Expected benefit:

- Faster service startup.

- 
- Cheaper incremental updates (only rebuild when the corpus changes significantly).

---

## 6.4 Use more advanced FAISS index types for large scale

Problem:

- A flat L2 index (IndexFlatL2) is simple and exact but not necessarily efficient for very large datasets.

Idea:

- For large collections, migrate to IVF-based or HNSW-based indexes:
  - IndexIVFFlat with a trained quantizer.
  - Potentially HNSW for better recall/speed trade-offs.

Expected benefit:

- Faster queries and lower memory usage at the cost of slightly approximate results, which is usually acceptable in semantic search.
- 

## 6.5 Warm-up and prefetching

Problem:

- First requests suffer from model loading and potential caching overhead.

Idea:

- On application startup:
    - Instantiate EmbeddingGenerator once.
    - Run a dummy call to `_encode_text` as a warm-up.
  - This shifts the cost to startup time instead of the first user-visible request.
- 

## 7. Next Steps

---

For later parts of A5 and future work, we plan to:

1. Implement at least one of the proposed optimizations in the actual codebase (for example, batching in generating `e_embeddings` and persisting FAISS indices).
2. Use Locust in A5 Part 2 to measure:
  - End-to-end latency for search APIs.
  - The effect of caching and warm-up on P95/P99 latency.
3. In A5 Part 3, design a scaling architecture that:
  - Separates embedding computation from query serving.
  - Allows horizontal scaling of the search service.

- Uses offline jobs for index building and refreshing.
- 

## 8. Files Included

---

- a5/profile\_a5\_part1.py
  - Profiling harness for the five target functions.
- a5/profiling\_results/
  - One .prof file per function (latest run):
    - profile\_encode\_text\_... .prof
    - profile\_generate\_embeddings\_... .prof
    - profile\_build\_index\_... .prof
    - profile\_search\_text\_... .prof
    - profile\_search\_category\_... .prof